




# GROUP 20 – REINFORCEMENT LEARNING






# Reinforcement Learning

- Supervised Learning; requires labelled data, can classify or regress
- Unsupervised Learning; unlabeled data, clustering and dimensionality reduction
- 3<sup>rd</sup> branch of machine learning; **Reinforcement Learning**



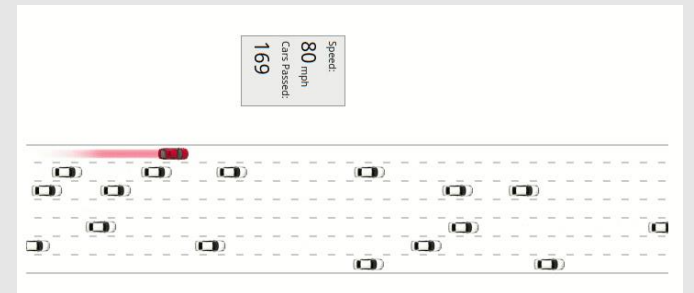


# Reinforcement Learning – Introduction

- Reinforcement Learning (RL):
  - Agent conducts action, is rewarded based on action
  - Learning via Trial and Error
  - Markov Decision Process (MDP) the bedrock of RL
- Important characteristics of RL
  - No supervisor, only reward signal
  - Process involves sequential decision making
  - Agents actions determine subsequent data
  - The agent must complete an objective but the actions to complete this are unknown
- Goal of RL
  - Take actions that maximize cumulative reward




# Reinforcement Learning – Applications and Success Stories

- AlphaGo; Google DeepMind AI defeated South Korean professional player Lee Sedol 4 matches to 1 in traditional “Go” game
- Robotics; find best combination of signals to complete a defined set of motions
- Autonomous driving; voyage deep drive, AWS deep racer, deep traffic




# Reinforcement Learning – Important Terms

- **Agent:** an entity which performs actions in an environment to gain some reward
- **Environment:** a scenario that an agent must face
- **Reward:** a return given to an agent when a specific task is performed
- **State:** the current condition returned by the environment
- **Policy:** strategy applied by agent to decide the next action based on current state
- **Value:** expected long term reward (i.e. collection of rewards) as compared to the short-term reward



# Markov Decision Process

- MDPs have a finite set of states ( $S$ ), a set of actions ( $A$ ), and a set of Rewards ( $R$ )



- The process of receiving a reward is an arbitrary function that maps state action pairs to rewards
$$f(S_t, A_t) = R_{t+1}$$
- The set of sequential processes that represent the MDP can be represented as  $S_t, A_t, R_{t+1}, \dots, S_n, A_n, R_{n+1}$
- Since our goal is to maximize the reward, we need to find a sequence of state, action pairs that achieve maximum cumulative R



# Solving MDPs with Q Learning – Cumulative Reward

- The goal of an agent is to maximize cumulative rewards; need to aggregate discrete R terms
- Can do this with the concept of expected return:

$$G_t = R_{t+1} + R_{t+2} + \dots R_T$$

- Where T is the final time step, in an episodic task
- For continuous tasks  $T = \infty$ , need to consider discounting
- Discount rate  $\gamma$  is a number between 0 and 1, used to determine the present value of future rewards

$$\begin{aligned} G_t &= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \\ &= \sum \gamma^k R_{t+k+1} \end{aligned}$$

- The discounted return ( $G_T$ ) makes it so that immediate rewards are worth more than future rewards, because future rewards are more heavily discounted



# Solving MDPs with Q Learning – Policies and Value Functions

## Policies

- What is the probability that an agent will select a specific action from a specific state?

## Value Functions

- How beneficial is being in a specific state or taking a specific action for the agent?





# Solving MDPs with Q Learning – Policies and Value Functions

## Policies

- If an agent follows a given policy ( $\pi$ ) at time  $t$  then  $\pi(a|s)$  is the probability that the agent will take action  $a$  in state  $s$  (or  $A_t = a$  &  $S_t = s$ )
- At time  $t$ , following policy  $\pi$ , the probability of an agent taking an action  $a$  in state  $s$  is  $\pi(a|s)$

## Value Functions

- Q learning uses action-value function ( $q_\pi$ ); gives the value of an action under policy  $\pi$
- The value of action  $a$ , in state  $s$ , under policy  $\pi$  is the expected return from starting from state  $s$  at time  $t$ , taking action  $a$ , and following policy  $\pi$
- Mathematically

$$q_\pi(s, a) = E[G_t | S_t = s, A_t = a]$$

Q-Function

Q-Value



# Solving MDPs with Q Learning – Optimal Action Value Function

- Optimal Q function ( $q_*$ ) gives the largest expected return achievable when following  $\pi$  for all possible state action pairs and can be defined as

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$

- For Q Learning,  $q_*$  must satisfy the Bellman Optimality equation

$$q_*(s, a) = E[R_{t+1} + \gamma \max_{a'} q_*(s', a')]$$

- The Q learning algorithm iteratively updates the Q values for a state-action pair using the Bellman equation until the Q function converges to  $q_*$ ; called value iteration
- Q values are stored in Q table, agent will take action that leads to highest Q value in a given state

	Action 1	Action 2	Action 3	Action 4
State 1	Q(1, 1)	Q(1, 2)	Q(1, 3)	Q(1, 4)
State 2	Q(2, 1)	Q(2, 2)	Q(2, 3)	Q(2, 4)
State 3	Q(3, 1)	Q(3, 2)	Q(3, 3)	Q(3, 4)



# Solving MDPs with Q Learning – Exploitation and Exploration

	Action 1	Action 2	Action 3	Action 4
State 1	0	0	0	0
State 2	0	0	0	0
State 3	0	0	0	0

?

- The agent must weigh **exploitation/exploration** tradeoff
  - Exploit what it already knows vs. explore rewards from untested actions
  - Can use the following to address exploitation exploration tradeoff
    - Random exploration
    - Greedy exploitation
    - $\epsilon$  - Greedy strategy



# Solving MDPs with Q Learning – Developing a Full Q Table

- What happens if agent is unable to find optimal Q values for each state action pair on first try?
- Do not want agent to forget everything from previous iteration if new iteration is started
- Objective is to develop a Q table with optimal Q values for each state action pair based on many iterations of agent in state action pairs (called episodes)
- Can be achieved with the **learning rate** ( $\alpha, \in [0, 1]$ )

$$q_{new}(s, a) = (1 - \alpha)q(s, a) + \alpha \left( R_{t+1} + \gamma \max_{a'} q(s', a') \right)$$

New table value

Old table value

New calculated value ( $q_*$ )

Episode 1

	Action 1
State 1	0



Episode 2

	Action 1
State 1	-0.5



Episode 3

	Action 1
State 1	-0.25



# The Q Learning Algorithm

Initialize Q  
table

Select  
Initial State

Perform an  
Action

Evaluate  
Q value  
and  
Update Q  
Table

Return to 2  
if goal is  
reached,  
else return  
to 3



# Q Learning Example – The Frozen Lake

- Open AI gym – Frozen Lake
- Frozen Lake game; Hiker crossing frozen lake
- Actions that lead to holes in the ice receive no reward
- Actions that lead to the party receive positive reward
- Over many iterations the actions that receive negative rewards will be filtered out while actions that receive positive rewards will remain



# Q Learning Example – The Frozen Lake

- Agent can move left, right, up, or down (4 actions)
- Agent can walk on 11 frozen lake tiles, 4 hole tiles, or 1 party tile (16 states)
- Walking to the party tile yields a reward of 1 while all other tiles yield reward of 0



# Q Learning Example – Define Helper Functions

```
In [3]: import gym
import numpy as np
import matplotlib.pyplot as plt
import time, pickle, os
from statistics import mean
%matplotlib inline
```

Library imports

```
def choose_action(state):
    action=0
    if np.random.uniform(0, 1) < epsilon:
        action = env.action_space.sample()
    else:
        action = np.argmax(Q[state, :])
    return action
```

Function which chooses an action based on a state

```
def learn(state, state2, reward, action):
    predict = Q[state, action]
    target = reward + gamma * np.max(Q[state2, :])
    Q[state, action] = Q[state, action] + lr_rate * (target - predict)
```

Function which Updates Q table





# Q Learning Example – Initialize Environment

In [5]: `env = gym.make('FrozenLake-v0')`

`epsilon = 0.9`  
`total_episodes = 2000`  
`max_steps = 100`

`lr_rate = 0.81`  
`gamma = 0.96`

`Q = np.zeros((env.observation_space.n, env.action_space.n))`  
`scores = []`

Initializes game environment

Defines q learning parameters

Initializes Q table



# Q Learning Example – Running the Episodes

```
In [6]: for episode in range(total_episodes):
        state = env.reset()
        t = 0

        while t < max_steps:
            env.render()

            action = choose_action(state)

            state2, reward, done, info = env.step(action)

            learn(state, state2, reward, action)

            state = state2

            t += 1

            print(episode)

            if done:
                break

            time.sleep(0.1)

        score = np.round(np.sum(Q/np.max(Q)*100))
        scores.append(score)

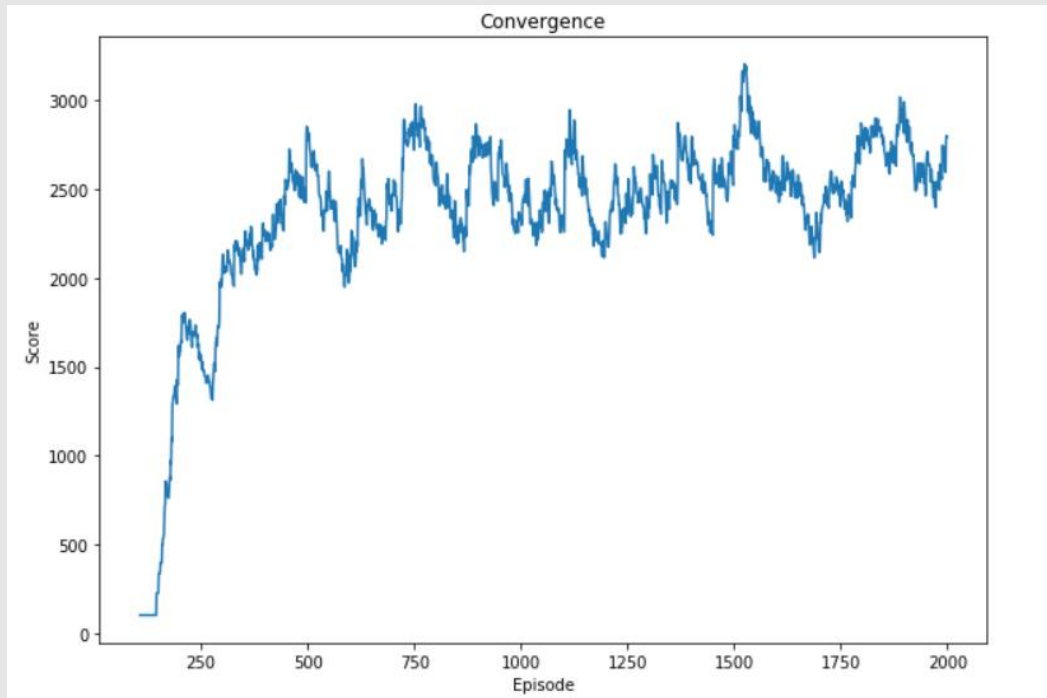
    print(Q)
```

Iterating through episodes, our agent takes actions, receives rewards, calculates q values and updates the q table based on the actions it has taken



# Q Learning Example – Convergence

Q Score






Final Q matrix

```
[[6.53118287e-01 6.76787165e-01 6.51955832e-01 6.62853501e-01]
 [6.56569162e-01 5.35738000e-01 5.72516132e-01 6.60733902e-01]
 [6.85009109e-01 6.86108158e-01 6.66498436e-01 6.80596059e-01]
 [5.57194560e-01 6.56040256e-01 6.47382293e-01 6.82087373e-01]
 [6.69494180e-01 6.34166430e-01 1.40381227e-01 6.53596992e-01]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [8.00347795e-01 3.09376784e-05 7.51117493e-01 1.30250050e-01]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [2.13311240e-02 7.83891691e-01 5.98714178e-01 7.12689471e-01]
 [7.14559204e-01 8.31028819e-01 8.86971198e-01 1.45849926e-01]
 [8.18272147e-01 9.26745089e-01 7.83663317e-01 1.52822043e-01]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]
 [4.87152449e-03 9.13074747e-01 9.34373235e-01 8.23645602e-01]
 [9.18358844e-01 9.07034013e-01 8.56813866e-01 9.95949472e-01]
 [0.00000000e+00 0.00000000e+00 0.00000000e+00 0.00000000e+00]]
```



# Q Learning Example – Hyperparameter Tuning

- Tuned hyper parameter to get agent to complete maze in fewest steps
- Hyper Parameter Tuning yielded an optimal  $\epsilon = 0.9$ ,  $\alpha = 0.3$ , and  $\gamma = 0.2$



Thanks!

